


ai-cv-screener — from PDF to a graded answer

Thirty synthetic CVs, read into Postgres and questioned in natural language. The eval suite checks both halves: what was read, and what was answered.

1 – Build the corpus · offline, once



— spends an LLM call - - - eval suite - - - file on disk — request or data flow

WHAT EACH STEP IS DOING

make generate

Writes the corpus **and** the answer key in one pass — every fact in a CV was decided before it was rendered, which is what makes grading possible later.

make ingest

One `POST /ingest`. Reads every PDF in `data/cvs`, then rebuilds the tables from scratch so a re-run never leaves half of an old corpus behind.

PDF → text

`pypdf` keeps a section's lines together; `pdfplumber` interleaves the columns but puts the spaces in. Both run, and each is used for what it is good at.

Extract fields

Ten CVs per call — the instruction block is the only part a batch shares, so batching past ten saves almost nothing and mixes people up more often.

data/extractions/

Cached against a hash of the CV text. An interrupted ingest never re-buys the tokens it already spent, and a changed CV invalidates itself.

Chunk + embed

900-character chunks with 150 of overlap, embedded locally with `bge-small` (384-d). Search costs nothing and needs no key.

POST /chat

One HTTP request, held open. Everything the answer needs — route, sources, then the text itself — comes back down it as Server-Sent Events.

QueryRouter

Matches the question against the words already in the corpus first, and only spends an LLM call when that fails. Most questions never reach the model twice.

The three routes

Vector search alone cannot answer *who has experience with Python?* — it returns the nearest eight chunks, not every match. Counting questions go to SQL; profile goes to one person's chunks; the vague ones stay semantic.

AnswerGenerator

Answers only from the CVs it was handed, cites the file behind each claim, and streams tokens as they arrive so the UI can start typing.

make eval-extraction

Scores what ingestion read out of every CV against the answer key — years, role, skills, institutions. Free, needs nothing running, and it caught the two columns fusing into `StackPython`.

make evals

One question of each shape, graded on the answer **and** on the route that produced it. The questions are derived from the answer key, so regenerating the corpus does not turn the suite into stale failures.